

DATE - 28/10/2023

PHASE - IV

TEAM ID - 696

PROJECT TITLE - House Pricing forecasting using ML



Importing Dependencies

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
import matplotlib.pyplot as plt
import tkinter as tk
import random
import requests
import scipy
```

Loading Dataset

```
In [3]: dataset = pd.read_csv("USA_Housing.csv")
```

Data importing

In [4]:

dataset

Out[4]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry / 674\nLaurabury, 370
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Vie Suite 079\nL: Kathleen, C.
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizab Stravenue\nDanielto' WI 0648
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO 446
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nF AE 090
...	
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\nF AP 30153-76
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 9258, E 8489\nAPO AA 429 30
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy Garc Suite 076\nJoshuala VA 0
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nFPO 730
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridg Apt. 509\nEast Hc NV

5000 rows × 7 columns

In [127]:

dataset.head(6)

Out[127]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386
5	80175.754159	4.988408	6.104512	4.04	26748.428425	1.068138e+06	06039 Jennifer Islands Apt. 443\nTracyport, KS...

In [6]:

dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                     5000 non-null   float64
1   Avg. Area House Age                  5000 non-null   float64
2   Avg. Area Number of Rooms            5000 non-null   float64
3   Avg. Area Number of Bedrooms         5000 non-null   float64
4   Area Population                      5000 non-null   float64
5   Price                               5000 non-null   float64
6   Address                             5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

In [7]:

dataset.describe()

Out[7]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
50%	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
75%	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
max	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06

In [8]: dataset.columns

Out[8]: Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms', 'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'], dtype='object')

```
In [163]: import pandas as pd
first_row = dataset.iloc[0]
specific_row = dataset.iloc[5]
print("First row:")
print(first_row)

print("\nSpecific row:")
print(specific_row)
```


```
First row:
Avg. Area Income      79545.4585
74
Avg. Area House Age    5.6828
61
Avg. Area Number of Rooms      7.0091
88
Avg. Area Number of Bedrooms    4.
09
Area Population      23086.8005
03
Price      1059033.557
87
Address      208 Michael Ferry Apt. 674\nLaurabury, NE 370
1...
Name: 0, dtype: object

Specific row:
Avg. Area Income      80175.7541
59
Avg. Area House Age    4.9884
08
Avg. Area Number of Rooms      6.1045
12
Avg. Area Number of Bedrooms    4.
04
Area Population      26748.4284
25
Price      1068138.0743
94
Address      06039 Jennifer Islands Apt. 443\nTracyport, K
S...
Name: 5, dtype: object
```

Pre-Processing and Visualisation of Data


```
In [9]: sns.histplot(dataset, x='Price', bins=50, color='r')
```

```
Out[9]: <Axes: xlabel='Price', ylabel='Count'>
```




```
In [10]: sns.histplot(dataset, y='Avg. Area Income', bins=50, color='b')
```

```
Out[10]: <Axes: xlabel='Count', ylabel='Avg. Area Income'>
```




```
In [12]: import pandas as pd
import seaborn as sns
dataset = pd.read_csv("USA_Housing.csv")
sns.barplot(x='Price', y='Avg. Area Income', data=dataset, palette='plasma')
plt.show()
```



```
In [11]: import seaborn as sns
import matplotlib.pyplot as plt

sns.boxplot(data=dataset, x='Price', palette='Blues')
plt.xlabel('Price')
plt.title('Boxplot of Price')


plt.show()
```



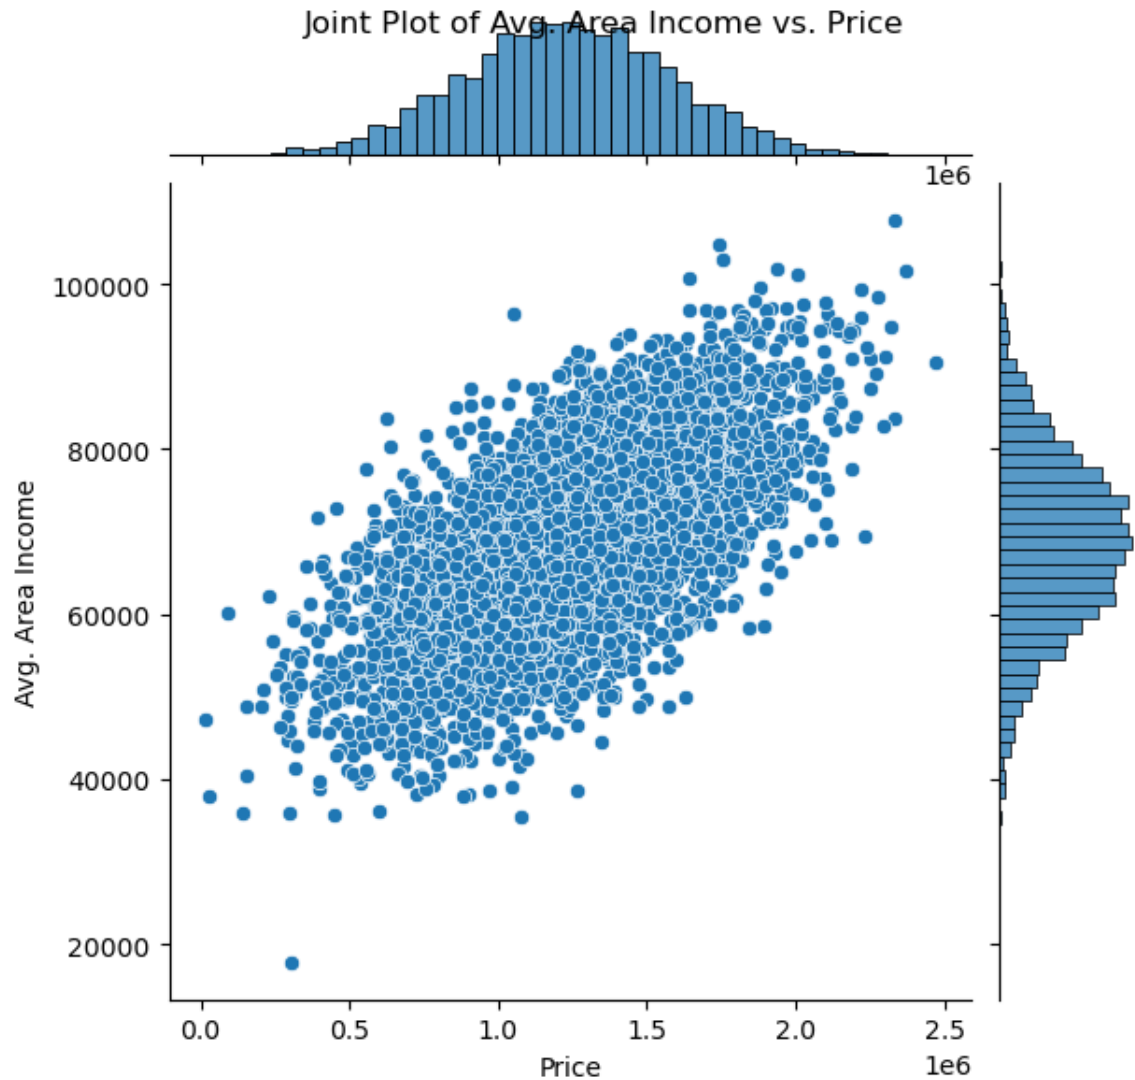
```
In [13]: import pandas as pd
import seaborn as sns
dataset = pd.read_csv("USA_Housing.csv")
import matplotlib.pyplot as plt

sns.jointplot(data=dataset, x='Price', y='Avg. Area House Age', kind='hex')
plt.xlabel('Price')
plt.ylabel('Avg. Area House Age')
plt.suptitle('Joint Plot with Hexagonal Bins')

plt.show()
```

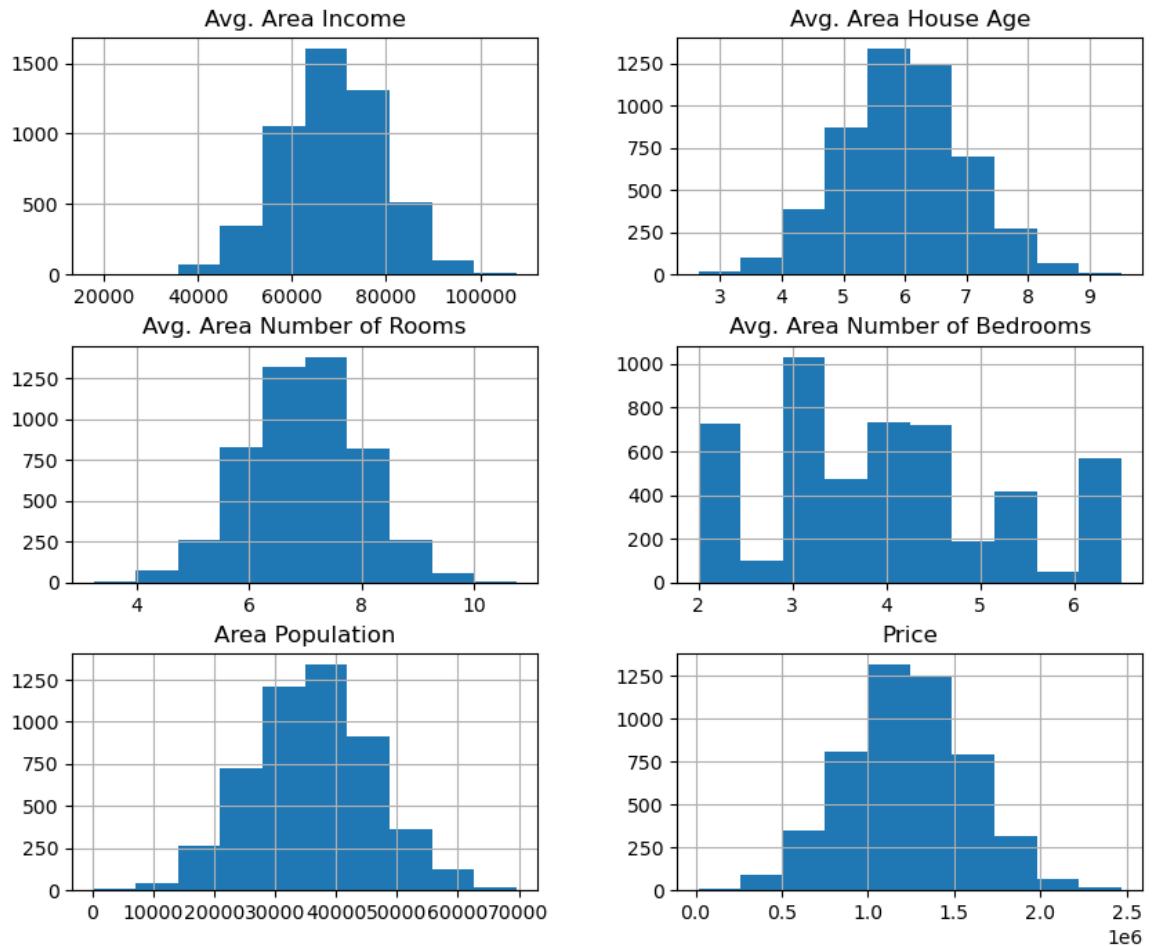



```
In [14]: import seaborn as sns
import matplotlib.pyplot as plt
sns.jointplot(data=dataset, x='Price', y='Avg. Area Income')
plt.xlabel('Price')
plt.ylabel('Avg. Area Income')
plt.suptitle('Joint Plot of Avg. Area Income vs. Price')
plt.show()
```



```
In [15]: dataset.hist(figsize=(10,8))
```

```
Out[15]: array([[<Axes: title={'center': 'Avg. Area Income'}>,  
  <Axes: title={'center': 'Avg. Area House Age'}>],  
  [<Axes: title={'center': 'Avg. Area Number of Rooms'}>,  
  <Axes: title={'center': 'Avg. Area Number of Bedrooms'}>],  
  [<Axes: title={'center': 'Area Population'}>,  
  <Axes: title={'center': 'Price'}>]], dtype=object)
```



Visualising Correlation

```
In [130]: dataset = dataset.select_dtypes(include=[np.number])  
correlation_matrix = dataset.corr()
```

```
In [17]: dataset.corr(numeric_only=True)
```

```
Out[17]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
Avg. Area Income	1.000000	-0.002007	-0.011032	0.019788	-0.016234	0.639734
Avg. Area House Age	-0.002007	1.000000	-0.009428	0.006149	-0.018743	0.452543
Avg. Area Number of Rooms	-0.011032	-0.009428	1.000000	0.462695	0.002040	0.335664
Avg. Area Number of Bedrooms	0.019788	0.006149	0.462695	1.000000	-0.022168	0.171071
Area Population	-0.016234	-0.018743	0.002040	-0.022168	1.000000	0.408556
Price	0.639734	0.452543	0.335664	0.171071	0.408556	1.000000

```
In [18]: import pandas as pd
dataset = pd.read_csv("USA_Housing.csv")
# Filter the DataFrame to include only numeric columns
numeric_columns = dataset.select_dtypes(include=['number'])
correlation_matrix = numeric_columns.corr()
print(correlation_matrix)
```

	Avg. Area Income	Avg. Area House Age \
Avg. Area Income	1.000000	-0.002007
Avg. Area House Age	-0.002007	1.000000
Avg. Area Number of Rooms	-0.011032	-0.009428
Avg. Area Number of Bedrooms	0.019788	0.006149
Area Population	-0.016234	-0.018743
Price	0.639734	0.452543

	Avg. Area Number of Rooms \
Avg. Area Income	-0.011032
Avg. Area House Age	-0.009428
Avg. Area Number of Rooms	1.000000
Avg. Area Number of Bedrooms	0.462695
Area Population	0.002040
Price	0.335664

	Avg. Area Number of Bedrooms	Area Population \
Avg. Area Income	0.019788	-0.016234
Avg. Area House Age	0.006149	-0.018743
Avg. Area Number of Rooms	0.462695	0.002040
Avg. Area Number of Bedrooms	1.000000	-0.022168
Area Population	-0.022168	1.000000
Price	0.171071	0.408556

	Price
Avg. Area Income	0.639734
Avg. Area House Age	0.452543
Avg. Area Number of Rooms	0.335664
Avg. Area Number of Bedrooms	0.171071
Area Population	0.408556
Price	1.000000

```
In [19]: plt.figure(figsize=(10,5))  
sns.heatmap(dataset.corr(numeric_only = True), annot=True)
```

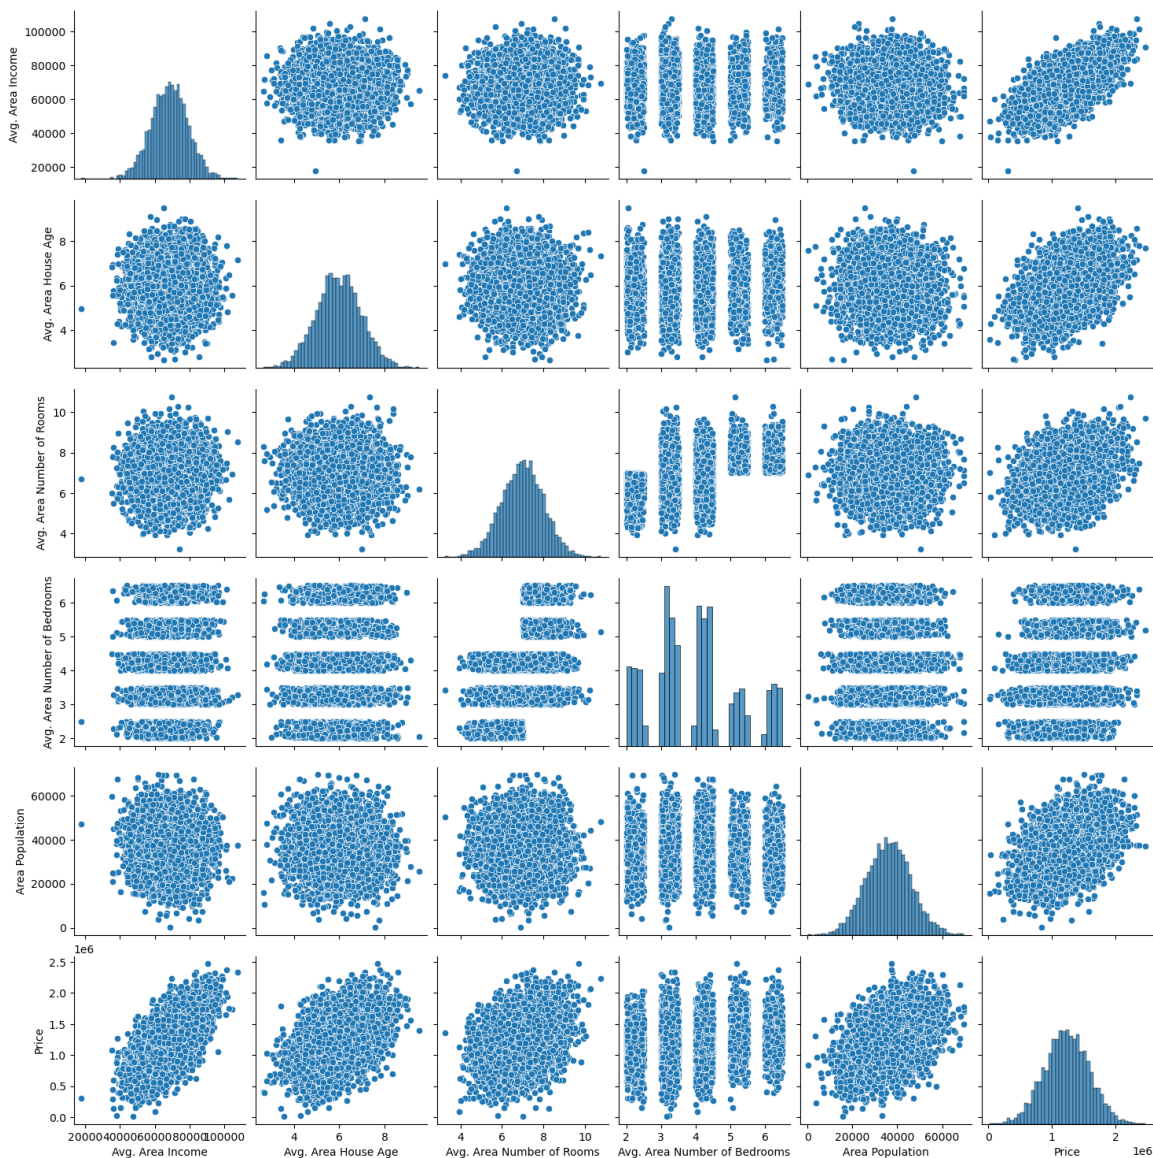
Out[19]: <Axes: >



```
In [69]: plt.figure(figsize=(12,8))
sns.pairplot(dataset)
```

Out[69]: <seaborn.axisgrid.PairGrid at 0x287aac16d90>

<Figure size 1200x800 with 0 Axes>



```
In [70]: dataset.isnull().sum()
```

```
Out[70]: Avg. Area Income      0
Avg. Area House Age      0
Avg. Area Number of Rooms 0
Avg. Area Number of Bedrooms 0
Area Population           0
Price                    0
Address                  0
dtype: int64
```

```
In [145]: ## importing modules

%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

In [149]: *## Dividing Data*

```
In [23]: X = dataset['Price']
Y = dataset[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms', 'Avg. Area Number of Bedrooms', 'Area Population']]
```

In [133]: *# TRAINING PHASE*

```
In [25]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_
```

In [26]: Y_train.head()

Out[26]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population
3413	69048.788093	6.619712	6.123813	4.33	36817.368760
1610	67866.899929	5.393978	9.359022	5.44	43122.574176
3459	56636.238191	5.497667	7.121872	6.10	47541.431763
4293	79310.361977	4.247434	7.518204	4.38	43982.188957
1039	72821.247664	6.480819	7.116655	5.33	40594.059297

In [27]: Y_train.shape

Out[27]: (4000, 5)

In [132]: *# TESTING PHASE*

In [28]: Y_test.head()

Out[28]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population
1718	66774.995817	5.717143	7.795215	4.32	36788.980327
2511	62184.539375	4.925758	7.427689	6.22	26008.309124
345	73643.057298	6.766853	8.337085	3.34	43152.139577
2521	61909.041438	6.228343	6.593138	4.29	28953.925377
54	72942.705059	4.786222	7.319886	6.41	24377.909049

In [29]: Y_test.shape

Out[29]: (1000, 5)

In [30]: *# Standardizing the data*

```
In [31]: from sklearn.preprocessing import StandardScaler
import numpy as np
sc = StandardScaler()
X_train_rescaled = X_train.to_numpy().reshape(-1, 1)
X_test_rescaled = X_test.to_numpy().reshape(-1, 1)
X_train_scaled = sc.fit_transform(X_train_rescaled)
X_test_scaled = sc.transform(X_test_rescaled)
```

```
In [32]: # Model Building and Evaluation
## Model 1 - Linear Regression
```

```
In [33]: from sklearn.linear_model import LinearRegression
model_lr = LinearRegression()
```

```
In [34]: model_lr.fit(X_train_scaled, Y_train)
```


```
Out[34]: ▾ LinearRegression
LinearRegression()
```

```
In [35]: ## Predicting Prices
```

```
In [36]: Prediction1 = model_lr.predict(X_test_scaled)
```

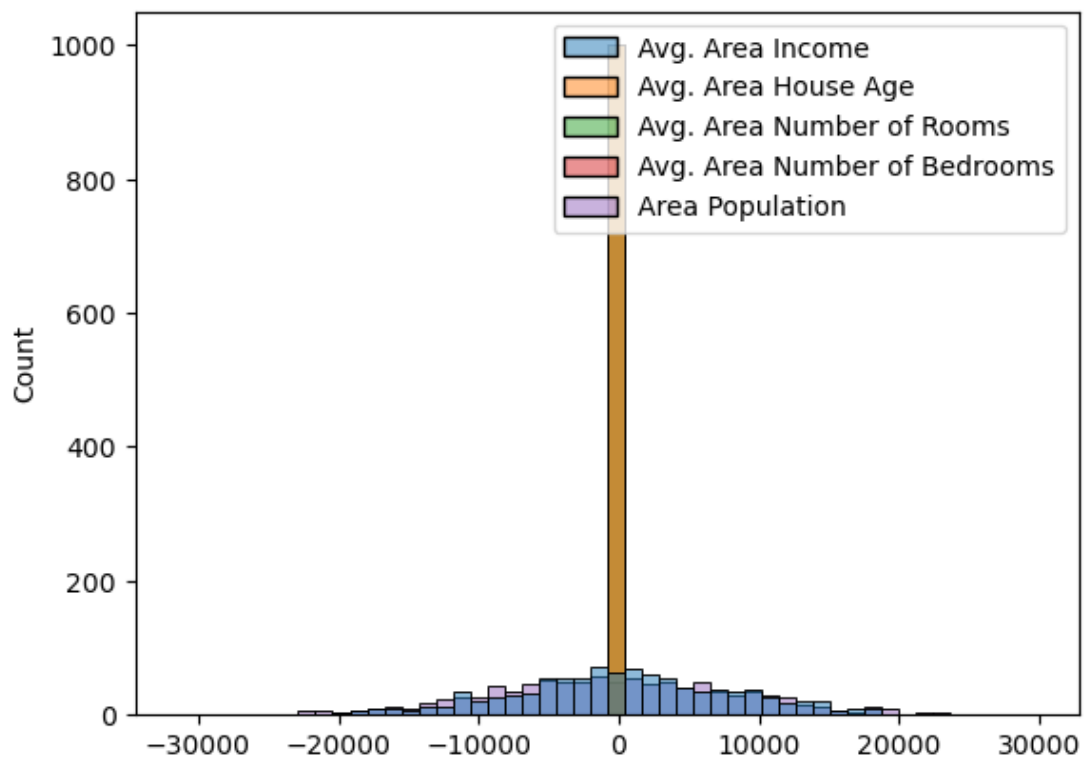
```
In [37]: ## Evaluation of Predicted Data
```

```
In [38]: import matplotlib.pyplot as plt
import numpy as np
if len(Y_test) == len(Prediction1):
    plt.figure(figsize=(12, 6))
    plt.plot(np.arange(len(Y_test)), Y_test, label='Actual Trend')
    plt.plot(np.arange(len(Prediction1)), Prediction1, label='Predicted Trend')
    plt.xlabel('Data')
    plt.ylabel('Trend')
    plt.legend()
    plt.title('Actual vs Predicted')
else:
    print("Lengths of Y_test and Prediction1 do not match.")
```



```
In [39]: sns.histplot((Y_test-Prediction1), bins=50,color='r')
```

Out[39]: <Axes: ylabel='Count'>




```
In [40]: print(r2_score(Y_test, Prediction1))  
print(mean_absolute_error(Y_test, Prediction1))  
print(mean_squared_error(Y_test, Prediction1))
```

```
0.1866240420321335  
2752.0362321666653  
30247095.384854764
```

```
In [76]: import numpy as np  
  
# Assuming X_train is a Pandas Series  
X_train_resaped = X_train.values.reshape(-1, 1)  
  
# Then, proceed with fitting the model using the reshaped data  
model_lr.fit(X_train_resaped, Y_train)
```


```
Out[76]: ▾ LinearRegression  
LinearRegression()
```

```
In [78]: print(model_lr.intercept_)
```

```
[4.47830789e+04 4.44480736e+00 5.81749836e+00 3.23285804e+00  
2.17417850e+04]
```

```
In [79]: sns.distplot(dataset["Price"])
```

```
Out[79]: <Axes: xlabel='Price', ylabel='Density'>
```



In [142]:


```
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
np.random.seed(0)
X = np.random.rand(100, 1)
y = 2 * X.squeeze() + np.random.randn(100)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# Creating and fitting the Linear Regression model
lm = LinearRegression()
lm.fit(X_train, y_train)

# Making predictions on the test set
predictions = lm.predict(X_test)

# Plotting the training data and the Linear regression line
plt.scatter(X, y, color='blue', label='Training Data')
plt.plot(X, lm.predict(X), color='red', label='Linear Regression Line')
plt.xlabel('X')
plt.ylabel('y')
plt.title('Linear Regression Example')
plt.legend()
plt.show()
```



```
In [144]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

np.random.seed(0)
X = np.random.rand(100, 1) # Feature
y = 2 * X.squeeze() + np.random.randn(100)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# Creating and fitting the Linear Regression model
lm = LinearRegression()
lm.fit(X_train, y_train)

# Creating a DataFrame with coefficients
coeff_df = pd.DataFrame(lm.coef_, columns=['Coefficient'])
print(coeff_df)
```

```
Coefficient
0      1.980518
```

```
In [86]: ## Model 2 - Support Vector Regressor
```

```
In [42]: model_svr = SVR()
```

```
In [43]: import numpy as np
Y_train_array = Y_train.to_numpy()
```

```
In [44]: from sklearn.multioutput import MultiOutputRegressor
from sklearn.svm import SVR
base_model = SVR(kernel='linear', C=1.0)

# Wrap it with MultiOutputRegressor
model_svr = MultiOutputRegressor(base_model)
model_svr.fit(X_train_scal, Y_train)
```


```
Out[44]: > MultiOutputRegressor
          > estimator: SVR
              > SVR
```

```
In [45]: ## Predicting Prices
```

```
In [46]: Prediction2 = model_svr.predict(X_test_scal)
```


```
In [47]: plt.figure(figsize=(12,6))
plt.plot(np.arange(len(Y_test)), Y_test, label='Actual Trend')
plt.plot(np.arange(len(Y_test)), Prediction2, label='Predicted Trend')
plt.xlabel('Data')
plt.ylabel('Trend')
plt.legend()
plt.title('Actual vs Predicted')
```

Out[47]: Text(0.5, 1.0, 'Actual vs Predicted')



```
In [48]: sns.histplot((Y_test-Prediction2), bins=50)
```

Out[48]: <Axes: ylabel='Count'>



```
In [49]: print(r2_score(Y_test, Prediction2))
print(mean_absolute_error(Y_test, Prediction2))
print(mean_squared_error(Y_test, Prediction2))
```

```
0.12092808540883056
3048.284745891538
37172622.89729237
```

```
In [157]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_


# Creating and fitting the Linear Regression model
lm = LinearRegression()
lm.fit(X_train, y_train)

# Creating a DataFrame with coefficients
coeff_df = pd.DataFrame(lm.coef_, columns=['Coefficient'])
print(coeff_df)
```

```
      Coefficient
0      1.980518
```

```
In [156]: import matplotlib.pyplot as plt

# Scatter plot of true test values vs. predicted values
plt.scatter(y_test, predictions)
plt.xlabel('True Values')
plt.ylabel('Predictions')
plt.title('True Values vs. Predicted Values')
plt.show()
```



In [50]: *## Model 3 - Lasso Regression*

In [51]: `model_lar = Lasso(alpha=1)`

In [52]: `model_lar.fit(X_train_scal,Y_train)`

Out[52]:

```

  ▾  Lasso
    Lasso(alpha=1)

```

In [53]: *## Predicting Prices*

In [54]: `Prediction3 = model_lar.predict(X_test_scal)`

In [55]: *## Evaluation of Predicted Data*


In [56]:

```

plt.figure(figsize=(12,6))
plt.plot(np.arange(len(Y_test)), Y_test, label='Actual Trend')
plt.plot(np.arange(len(Y_test)), Prediction1, label='Predicted Trend')
plt.xlabel('Data')
plt.ylabel('Trend')
plt.legend()
plt.title('Actual vs Predicted')


```

Out[56]: `Text(0.5, 1.0, 'Actual vs Predicted')`



In [57]: `sns.histplot((Y_test-Prediction3), bins=50)`

Out[57]: <Axes: ylabel='Count'>



In [58]: `print(r2_score(Y_test, Prediction2))`
`print(mean_absolute_error(Y_test, Prediction2))`
`print(mean_squared_error(Y_test, Prediction2))`

0.12092808540883056
 3048.284745891538
 37172622.89729237

In [59]: `## Model 4 - Random Forest Regressor`

In [60]: `model_rf = RandomForestRegressor(n_estimators=50)`

In [61]: `model_rf.fit(X_train_scal, Y_train)`

Out[61]:

RandomForestRegressor
RandomForestRegressor(n_estimators=50)


In [62]: `## Predicting Prices`

In [63]: `Prediction4 = model_rf.predict(X_test_scal)`

In [64]: `## Evaluation of Predicted Data`


```
In [65]: plt.figure(figsize=(12,6))
plt.plot(np.arange(len(Y_test)), Y_test, label='Actual Trend')
plt.plot(np.arange(len(Y_test)), Prediction4, label='Predicted Trend')
plt.xlabel('Data')
plt.ylabel('Trend')
plt.legend()
plt.title('Actual vs Predicted')
```

Out[65]: Text(0.5, 1.0, 'Actual vs Predicted')



```
In [66]: sns.histplot((Y_test-Prediction4), bins=50)
```

Out[66]: <Axes: ylabel='Count'>




```
In [67]: print(r2_score(Y_test, Prediction3))
print(mean_absolute_error(Y_test, Prediction2))
print(mean_squared_error(Y_test, Prediction2))
```

```
0.11034396748383535
3048.284745891538
37172622.89729237
```

```
In [ ]:
```

```
In [167]: income_value = dataset.at[0, 'Avg. Area Income']
population_value = dataset.at[2, 'Area Population']
price_value = dataset.at[3, 'Price']
rooms_value = dataset.at[1, 'Avg. Area Number of Rooms']
Bedroom_value = dataset.at[4, 'Avg. Area Number of Bedrooms']
```

```
In [164]: income_value
```

```
Out[164]: 79545.45857431678
```

```
In [165]: population_value
```

```
Out[165]: 36882.15939970458
```

```
In [96]: price_value
```

```
Out[96]: 1260616.8066294468
```

```
In [166]: rooms_value
```

```
Out[166]: 6.730821019094919
```

```
In [168]: Bedroom_value
```

```
Out[168]: 4.23
```

```
In [99]: import numpy as np
import pandas as pd # Use 'pandas' instead of 'panda'

# Load your dataset into a Pandas DataFrame
dataset = pd.read_csv("USA_Housing.csv")

# Calculate the average
average_price = dataset['Price'].mean()
print("Average Price:", average_price)
```

```
Average Price: 1232072.654142357
```

```
In [100]: from sklearn.metrics import accuracy_score
```

```
In [102]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score

# Load your dataset into a Pandas DataFrame
dataset = pd.read_csv('USA_Housing.csv')

# Assuming 'Price' is your target column
# Split the dataset into features (X) and the target variable (y)
X = dataset.drop(columns=['Price', 'Address'])
y = dataset['Price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

# Create and train a machine learning model (RandomForestRegressor) using X_train
model = RandomForestRegressor()
model.fit(X_train, y_train)

# Use the trained model to make predictions on the test data (X_test)
y_pred = model.predict(X_test)

# Evaluate the regression model using mean squared error and R-squared (R2) score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print or use the evaluation metrics
print("Mean Squared Error:", mse)
print("R-squared (R2) Score:", r2)
```

Mean Squared Error: 15010250912.2045

R-squared (R2) Score: 0.8828507619666764


In [136]: !pip install basemap

```
Requirement already satisfied: basemap in c:\users\vijayraj r\anaconda3\lib\site-packages (1.3.8)
Requirement already satisfied: basemap-data<1.4,>=1.3.2 in c:\users\vijayraj r\anaconda3\lib\site-packages (from basemap) (1.3.2)
Requirement already satisfied: pyshp<2.4,>=1.2 in c:\users\vijayraj r\anaconda3\lib\site-packages (from basemap) (2.3.1)
Requirement already satisfied: matplotlib<3.8,>=1.5 in c:\users\vijayraj r\anaconda3\lib\site-packages (from basemap) (3.7.2)
Requirement already satisfied: pyproj<3.7.0,>=1.9.3 in c:\users\vijayraj r\anaconda3\lib\site-packages (from basemap) (3.6.1)
Requirement already satisfied: numpy<1.26,>=1.21 in c:\users\vijayraj r\anaconda3\lib\site-packages (from basemap) (1.24.3)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\vijayraj r\anaconda3\lib\site-packages (from matplotlib<3.8,>=1.5->basemap) (1.0.5)
Requirement already satisfied: cycler>=0.10 in c:\users\vijayraj r\anaconda3\lib\site-packages (from matplotlib<3.8,>=1.5->basemap) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\vijayraj r\anaconda3\lib\site-packages (from matplotlib<3.8,>=1.5->basemap) (4.25.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\vijayraj r\anaconda3\lib\site-packages (from matplotlib<3.8,>=1.5->basemap) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\vijayraj r\anaconda3\lib\site-packages (from matplotlib<3.8,>=1.5->basemap) (23.1)
Requirement already satisfied: pillow>=6.2.0 in c:\users\vijayraj r\anaconda3\lib\site-packages (from matplotlib<3.8,>=1.5->basemap) (9.4.0)
Requirement already satisfied: pyparsing<3.1,>=2.3.1 in c:\users\vijayraj r\anaconda3\lib\site-packages (from matplotlib<3.8,>=1.5->basemap) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\vijayraj r\anaconda3\lib\site-packages (from matplotlib<3.8,>=1.5->basemap) (2.8.2)
Requirement already satisfied: certifi in c:\users\vijayraj r\anaconda3\lib\site-packages (from pyproj<3.7.0,>=1.9.3->basemap) (2023.7.22)
Requirement already satisfied: six>=1.5 in c:\users\vijayraj r\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib<3.8,>=1.5->basemap) (1.16.0)
```

In [104]: **from** mpl_toolkits.basemap **import** Basemap

```
In [135]: import matplotlib.pyplot as plt

plt.figure(figsize=(15, 5))
bar_graph = dataset.groupby('Price')['Avg. Area Income'].mean()
bar_graph.plot(kind='bar', color='Green')
plt.xlabel('Price')
plt.ylabel('Mean Avg. Area Income')
plt.title('Mean Avg. Area Income by Price')
plt.show()
```



```
In [113]: len(X_test)
```

```
Out[113]: 1000
```

```
In [114]: len(Y_test)
```


```
Out[114]: 1000
```

```
In [158]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

dataset = pd.read_csv("USA_Housing.csv")

# Extracting the data
X = dataset['Price']
Y = dataset['Avg. Area Income']

sns.barplot(x=X, y=Y, data=dataset, palette='viridis')
plt.show()
```



In []: